

Gaussian Elimination High-Performance Implementation

Mohamed Imad Eddine Ghodbane
May, 2024



CONTENT

- 01** PROJECT OVERVIEW
- 02** PLUQ FACTORIZATION
- 03** MODULAR ARITHMETIC USING SIMD
- 04** SIMD IMPLEMENTATION OF PLUQ
- 05** CROUT METHOD
- 06** SIMD SCALAR PRODUCT
- 07** SIMD IMPLEMENTATION OF CROUT METHOD
- 08** RESULTS

PROJECT OVERVIEW

Primary Objective

- Develop a high-performance implementation of Gaussian elimination.
- Focus on exact linear algebra over finite fields ($F_p = \mathbb{Z}/p\mathbb{Z}$) with a prime number stored on 30 bits.

Existing Libraries

- FFLAS-FFPACK, Flint, NTL.
- High performance for large matrices and small prime fields.

Areas for Improvement

- Intermediate matrix dimensions (hundreds to thousands).
- Larger prime numbers.

PLUQ FACTORIZATION

$$A = PLUQ$$

Input

- Matrix A of size $m \times n$ with entries in the finite field F_p .

Output



- LU decomposition of A , with permutation matrices P and Q .

Key Operations

- Swap rows and columns.
- Zero out elements below the pivot (Row Reduction).

PLUQ FACTORIZATION

Data Representation



- P and Q are represented as vectors.
- Compact storage for L and U .

PLUQ FACTORIZATION

Column Transposition

	0				

Q

1
2
3
4
5
6

- Choose the next first non-zero entry in the row.
- Column transposition to move the non-zero entry to the pivot position.

PLUQ FACTORIZATION

Column Transposition

		0			

Q

1
3
2
4
5
6

- Choose the next first non-zero entry in the row.
- Column transposition to move the non-zero entry to the pivot position.

PLUQ FACTORIZATION

Row Rotation

	0	0	0	0	0

P

1
2
3
4
5
6

- Perform row rotation to move the zero row to the last row position.

PLUQ FACTORIZATION

Row Rotation

	0	0	0	0	0

P



1
3
4
5
6
2

- Perform row rotation to move the zero row to the last row position.

PLUQ FACTORIZATION

Row Reduction

$$l_{ij} = u_{ii}^{-1} \cdot a_{ij}$$





- Compute the inverse of the pivot modulo p .
- Multiply the pivot inverse by the elements of the pivot row.
- Use the updated pivot row to eliminate the elements below the pivot in the current column.

PLUQ FACTORIZATION

Row Reduction

$$a_{ij} = a_{ij} - l_{ik} \cdot u_{kj}$$



- Compute the inverse of the pivot modulo p .
- Multiply the pivot inverse by the elements of the pivot row.
- Use the updated pivot row to eliminate the elements below the pivot in the current column.

PLUQ FACTORIZATION

Row Reduction

$$l_{ij} = u_{ii}^{-1} \cdot a_{ij}$$

$$a_{ij} = a_{ij} - l_{ik} \cdot u_{kj}$$

- Compute the inverse of the pivot modulo p .
- Multiply the pivot inverse by the elements of the pivot row.
- Use the updated pivot row to eliminate the elements below the pivot in the current column.

PLUQ FACTORIZATION

Row Reduction

$$l_{ij} = u_{ii}^{-1} \cdot a_{ij}$$

$$a_{ij} = a_{ij} - l_{ik} \cdot u_{kj}$$

- Compute the inverse of the pivot modulo p .
- Multiply the pivot inverse by the elements of the pivot row.
- Use the updated pivot row to eliminate the elements below the pivot in the current column.

PLUQ FACTORIZATION

Row Reduction

$$l_{ij} = u_{ii}^{-1} \cdot a_{ij}$$

$$a_{ij} = a_{ij} - l_{ik} \cdot u_{kj}$$

- Compute the inverse of the pivot modulo p .
- Multiply the pivot inverse by the elements of the pivot row.
- Use the updated pivot row to eliminate the elements below the pivot in the current column.

PLUQ FACTORIZATION

Row Reduction

$$l_{ij} = u_{ii}^{-1} \cdot a_{ij}$$


$$a_{ij} = a_{ij} - l_{ik} \cdot u_{kj}$$

- Compute the inverse of the pivot modulo p .
- Multiply the pivot inverse by the elements of the pivot row.
- Use the updated pivot row to eliminate the elements below the pivot in the current column.

PLUQ FACTORIZATION


Negligible Operation

$$l_{ij} = u_{ii}^{-1} \cdot a_{ij}$$



15 Operations

$$a_{ij} = a_{ij} - l_{ik} \cdot u_{kj}$$



110 Operations

PLUQ FACTORIZATION

One Iteration

```
/*  
    Find Pivot  
*/  
  
// Row Reduction  
for (int k = matrixRank + 1; k < m; k++) {  
    A->data[k * n + matrixRank] = mult(  
        A->data[k * n + matrixRank], inv, p);  
    for (int j = matrixRank + 1; j < n; j++)  
        A->data[k * n + j] = sub(  
            A->data[k * n + j],  
            mult(A->data[k * n + matrixRank],  
                A->data[matrixRank * n + j], p), p);  
}
```

MODULAR ARITHMETIC USING SIMD

Subtraction

```
int sub(int a, int b, int p) {  
    int r = a - b;  
    return r < 0 ? r + p : r;  
}
```

Multiplication

```
int mult(int a, int b, int p) {  
    long long r = (long long) a * b;  
    return r % p;  
}
```

MODULAR ARITHMETIC USING SIMD

Subtraction


```
__m128i res = _mm_sub_epi32(a, b);  
__m128i mask = _mm_cmplt_epi32(res,  
                                _mm_setzero_si128());  
__m128i ar = _mm_add_epi32(res, p);  
res = _mm_blendv_epi8(res, ar, mask);
```

a_1	a_2	a_3	a_4
-			
b_1	b_2	b_3	b_4
=			
$a_1 - b_1$	$a_2 - b_2$	$a_3 - b_3$	$a_4 - b_4$

MODULAR ARITHMETIC USING SIMD

Subtraction


```
__m128i res = _mm_sub_epi32(a, b);  
__m128i mask = _mm_cmplt_epi32(res,  
    _mm_setzero_si128());  
__m128i ar = _mm_add_epi32(res, p);  
res = _mm_blendv_epi8(res, ar, mask);
```



MODULAR ARITHMETIC USING SIMD

Subtraction


```
__m128i res = _mm_sub_epi32(a, b);  
__m128i mask = _mm_cmplt_epi32(res,  
                                _mm_setzero_si128());  
__m128i ar = _mm_add_epi32(res, p);  
res = _mm_blendv_epi8(res, ar, mask);
```



MODULAR ARITHMETIC USING SIMD

Subtraction

```
__m128i res = _mm_sub_epi32(a, b);  
__m128i mask = _mm_cmplt_epi32(res,  
    _mm_setzero_si128());  
__m128i ar = _mm_add_epi32(res, p);  
res = _mm_blendv_epi8(res, ar, mask);
```



MODULAR ARITHMETIC USING SIMD

Multiplication

```
int mult(int a, int b, int p) {  
    long long r = (long long)a * b;  
    return r % p;  
}
```

There is no modulus
operation in AVX2 !!

Solution: compute the remainder of the division by p .

$$h = ab = cp + e$$

$$d = \frac{h}{p}$$


$$c = \lfloor d \rfloor$$

$$e = h - cp$$

MODULAR ARITHMETIC USING SIMD

Multiplication

```
__m256d h = _mm256_mul_pd(a, b);  
__m256d d = _mm256_mul_pd(h, u);  
__m256d c = _mm256_floor_pd(d);  
__m256d e = _mm256_fnmadd_pd(c, p, h);
```



MODULAR ARITHMETIC USING SIMD

Multiplication

```
__m256d h = _mm256_mul_pd(a, b);  
__m256d d = _mm256_mul_pd(h, u);  
__m256d c = _mm256_floor_pd(d);  
__m256d e = _mm256_fnmadd_pd(c, p, h);
```

$$\begin{array}{|c|c|c|c|} \hline h_1 & h_2 & h_3 & h_4 \\ \hline \end{array} * \begin{array}{|c|c|c|c|} \hline u_1 & u_2 & u_3 & u_4 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline d_1 & d_2 & d_3 & d_4 \\ \hline \end{array}$$

MODULAR ARITHMETIC USING SIMD

Multiplication

```
__m256d h = _mm256_mul_pd(a, b);  
__m256d d = _mm256_mul_pd(h, u);  
__m256d c = _mm256_floor_pd(d);  
__m256d e = _mm256_fnmadd_pd(c, p, h);
```

$$\begin{bmatrix} \lfloor d_1 \rfloor & \lfloor d_2 \rfloor & \lfloor d_3 \rfloor & \lfloor d_4 \rfloor \end{bmatrix} = \begin{bmatrix} c_1 & c_2 & c_3 & c_4 \end{bmatrix}$$

MODULAR ARITHMETIC USING SIMD

Multiplication

```
__m256d h = _mm256_mul_pd(a, b);  
__m256d d = _mm256_mul_pd(h, u);  
__m256d c = _mm256_floor_pd(d);  
__m256d e = _mm256_fnmadd_pd(c, p, h);
```

$$\begin{array}{|c|c|c|c|} \hline h_1 & h_2 & h_3 & h_4 \\ \hline \end{array} - \begin{array}{|c|c|c|c|} \hline c_1p & c_2p & c_3p & c_4p \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline e_1 & e_2 & e_3 & e_4 \\ \hline \end{array}$$

MODULAR ARITHMETIC USING SIMD

Multiplication

```
__m256d mul_mod_p(__m256d a, __m256d b, __m256d u, __m256d p) {  
    __m256d h = _mm256_mul_pd(a, b);  
    __m256d d = _mm256_mul_pd(h, u);  
    __m256d c = _mm256_floor_pd(d);  
    __m256d e = _mm256_fnmadd_pd(c, p, h);  
    return e;  
}
```

The Product exceeds the 52-bit length . This can lead to a loss of precision in the final result !!



MODULAR ARITHMETIC USING SIMD

Multiplication

```
__m256d mul_mod_p(__m256d a, __m256d b, __m256d u, __m256d p) {  
    __m256d h = _mm256_mul_pd(a, b);  
    __m256d l = _mm256_fmsub_pd(a, b, h);  
    __m256d d = _mm256_mul_pd(h, u);  
    __m256d c = _mm256_floor_pd(d);  
    __m256d b = _mm256_fnmadd_pd(c, p, h);  
    __m256d e = _mm256_add_pd(b, l);  
    __m256d t = _mm256_sub_pd(e, p);  
    e = _mm256_blendv_pd(t, e, t);  
    t = _mm256_add_pd(e, p);  
    return _mm256_blendv_pd(e, t, e);  
}
```

SIMD IMPLEMENTATION OF PLUQ

```
rows_elimination_avx2(int *A_data, int n, int matrixRank, int c, int p ,
__m256d vp, __m256d vu , __m128i vp_128, int k) {
    __m256d vc = _mm256_set1_pd(c);
    __m256d tmp;
    int i;
    for (i = matrixRank + 1; i + 3 < n; i += 4) {
        __m128i v1 = _mm_loadu_si128((__m128i *)&A_data[matrixRank*n+i]);
        __m128i v2 = _mm_loadu_si128((__m128i *)&A_data[k*n+i]);
        __m256d vDouble = _mm256_cvtepi32_pd(v1);
        tmp = mul_mod_p(vc, vDouble, vu, vp);
        __m128i resultInt = _mm256_cvttpd_epi32(tmp);
        __m128i result = sub_avx2(v2, resultInt, vp_128);
        _mm_storeu_si128((__m128i *)&A_data[k * n + i], result);
    }
    // loop handles elements that don't fit into chunks of 4
}
```

RESULTS


AVX2 vs Basic PLUQ

Sizes	100	300	500	1000	1200
Basic (ms)	0.69	18.70	86.57	694.06	1199.92
AVX2 (ms)	0.15	3.90	17.80	139.98	242.56
Speedup	4.60	4.79	4.86	4.95	4.95

PLUQ Speedups using AVX2 and 12 Bits Length Prime

RESULTS

AVX2 vs FLINT



RESULTS

AVX2 vs FFLAS-FFPACK

Sizes	100	300	500	1000	1200
FFLAS-FFPACK (ms)	0.34	4.43	15.01	62.32	88.28
AVX2 (ms)	0.15	3.90	17.80	139.98	242.56
Speedup	2.26	1.13	0.85	0.42	0.35

FFLAS-FFPACK Speedups using AVX2 and 12 Bits Length Prime

CROUT METHOD

Algorithm 2 Crout variant of PLUQ with lexicographic search and column rotations.

```
1:  $k \leftarrow 1$ 
2: for  $i = 1 \dots m$  do
3:    $A_{i,k..n} \leftarrow A_{i,k..n} - A_{i,1..k-1} \times A_{1..k-1,k..n}$ 
4:   if  $A_{i,k..n} = 0$  then
5:     Loop to next iteration
6:   end if
7:   Let  $A_{i,s}$  be the left-most nonzero element of row  $i$ .
8:    $A_{i+1..m,s} \leftarrow A_{i+1..m,s} - A_{i+1..m,1..k-1} \times A_{1..k-1,s}$ 
9:    $A_{i+1..m,s} \leftarrow A_{i+1..m,s} / A_{i,s}$ 
10:  Bring  $A_{*,s}$  to  $A_{*,k}$  by column rotation
11:  Bring  $A_{i,*}$  to  $A_{k,*}$  by row rotation
12:   $k \leftarrow k + 1$ 
13: end for
```

“Fast computation of the rank profile matrix and the generalized bruhat decomposition”
written by Jean-Guillaume Dumas, Clément Pernet, and Ziad Sultan.

CROUT METHOD

Key Operation

- Computes u_{ij} and l_{ij} using the multiplication equation $A=LU$ and the previously computed values of L and U .


$$A = \begin{pmatrix} a_{00} & a_{01} & a_{02} \\ a_{10} & a_{11} & a_{12} \\ a_{20} & a_{21} & a_{22} \end{pmatrix} = P \cdot \begin{pmatrix} 1 & 0 & 0 \\ l_{10} & 1 & 0 \\ l_{20} & l_{20} & 1 \end{pmatrix} \cdot \begin{pmatrix} u_{00} & u_{01} & u_{02} \\ 0 & u_{11} & u_{12} \\ 0 & 0 & u_{22} \end{pmatrix} \cdot Q$$

$$A = P \cdot \begin{pmatrix} u_{00} & u_{01} & u_{02} \\ l_{10}u_{00} & l_{10}u_{01} + u_{11} & l_{10}u_{02} + u_{12} \\ l_{20}u_{00} & l_{20}u_{01} + l_{21}u_{11} & l_{20}u_{02} + l_{21}u_{12} + u_{22} \end{pmatrix} \cdot Q$$

CROUT METHOD

Row Reduction


$$u_{ij} = a_{ij} - \sum_{k=0}^{j-1} l_{ik} \cdot u_{ki}$$



CROUT METHOD

Row Reduction

$$u_{ij} = a_{ij} - \sum_{k=0}^{j-1} l_{ik} \cdot u_{ki}$$




Green	Green	Green	Dark Green	Green	Green
Blue	Green	Green	Dark Green	Green	Green
Blue	Blue	Green	Gray	Light Gray	Light Gray
Blue	Blue	Light Gray	Light Gray	Light Gray	Light Gray
Blue	Blue	Light Gray	Light Gray	Light Gray	Light Gray
Blue	Blue	Light Gray	Light Gray	Light Gray	Light Gray

CROUT METHOD

Row Reduction

$$u_{ij} = a_{ij} - \sum_{k=0}^{j-1} l_{ik} \cdot u_{ki}$$




Green	Green	Green	Green	Dark Green	Green
Blue	Green	Green	Green	Dark Green	Green
Blue	Blue	Green	Green	Gray	Light Gray
Blue	Blue	Light Gray	Light Gray	Light Gray	Light Gray
Blue	Blue	Light Gray	Light Gray	Light Gray	Light Gray
Blue	Blue	Light Gray	Light Gray	Light Gray	Light Gray

CROUT METHOD

Row Reduction

$$u_{ij} = a_{ij} - \sum_{k=0}^{j-1} l_{ik} \cdot u_{ki}$$




Green	Green	Green	Green	Green	Dark Green
Blue	Green	Green	Green	Green	Dark Green
Blue	Blue	Green	Green	Green	Gray
Blue	Blue	Gray	Gray	Gray	Gray
Blue	Blue	Gray	Gray	Gray	Gray
Blue	Blue	Gray	Gray	Gray	Gray

CROUT METHOD

Row Reduction


$$l_{ij} = \left(a_{ij} - \sum_{k=0}^{i-1} l_{ik} \cdot u_{kj} \right) \cdot u_{jj}^{-1}$$



CROUT METHOD

Row Reduction


$$l_{ij} = \left(a_{ij} - \sum_{k=0}^{i-1} l_{ik} \cdot u_{kj} \right) \cdot u_{jj}^{-1}$$



CROUT METHOD

Row Reduction

$$l_{ij} = \left(a_{ij} - \sum_{k=0}^{i-1} l_{ik} \cdot u_{kj} \right) \cdot u_{jj}^{-1}$$



Light Green	Light Green	Dark Green	Light Green	Light Green	Light Green
Light Blue	Light Green	Dark Green	Light Green	Light Green	Light Green
Light Blue	Light Blue	Red	Light Green	Light Green	Light Green
Light Blue	Light Blue	Light Blue	Light Grey	Light Grey	Light Grey
Light Blue	Light Blue	Light Blue	Light Grey	Light Grey	Light Grey
Light Blue	Light Blue	Dark Grey	Light Grey	Light Grey	Light Grey

CROUT METHOD

Row Reduction

$$u_{ij} = a_{ij} - \sum_{k=0}^{j-1} l_{ik} \cdot u_{ki}$$
$$l_{ij} = \left(a_{ij} - \sum_{k=0}^{i-1} l_{ik} \cdot u_{kj} \right) \cdot u_{jj}^{-1}$$

CROUT METHOD

Row Reduction

$$u_{ij} = a_{ij} - \sum_{k=0}^{j-1} l_{ik} \cdot u_{ki}$$
$$l_{ij} = \left(a_{ij} - \sum_{k=0}^{i-1} l_{ik} \cdot u_{kj} \right) \cdot u_{jj}^{-1}$$

CROUT METHOD

Row Reduction

$$u_{ij} = a_{ij} - \sum_{k=0}^{j-1} l_{ik} \cdot u_{ki}$$
$$l_{ij} = \left(a_{ij} - \sum_{k=0}^{i-1} l_{ik} \cdot u_{kj} \right) \cdot u_{jj}^{-1}$$

CROUT METHOD

Row Reduction

$$u_{ij} = a_{ij} - \sum_{k=0}^{j-1} l_{ik} \cdot u_{ki}$$
$$l_{ij} = \left(a_{ij} - \sum_{k=0}^{i-1} l_{ik} \cdot u_{kj} \right) \cdot u_{jj}^{-1}$$

CROUT METHOD


Row Reduction

$$u_{ij} = a_{ij} - \sum_{k=0}^{j-1} l_{ik} \cdot u_{ki}$$
$$l_{ij} = \left(a_{ij} - \sum_{k=0}^{i-1} l_{ik} \cdot u_{kj} \right) \cdot u_{jj}^{-1}$$

CROUT METHOD

Row Reduction

$$u_{ij} = a_{ij} - \sum_{k=0}^{j-1} l_{ik} \cdot u_{ki}$$
$$l_{ij} = \left(a_{ij} - \sum_{k=0}^{i-1} l_{ik} \cdot u_{kj} \right) \cdot u_{jj}^{-1}$$




Crout's method uses 74 operations for modular reduction, while the basic PLUQ method uses 125.

SIMD SCALAR PRODUCT

Scalar Product


```
int sum = 0;
for (int i = 0; i < length; i++) {
    sum += a[i] * b[i];
}
return sum % p;
```



SIMD SCALAR PRODUCT


Scalar Product

```
int sum1 = 0;
int sum1 = 0;
int sum1 = 0;
int sum1 = 0;
for (int i = 0; i < length; i+=4) {
    sum1 += a[i] * b[i];
    sum2 += a[i+1] * b[i+1];
    sum3 += a[i+2] * b[i+2];
    sum4 += a[i+3] * b[i+3];
}
return (sum1+sum2+sum3+sum4) % p;
```



SIMD SCALAR PRODUCT

Scalar Product



$$result = (sum_1 + sum_2 + sum_3 + sum_4) \bmod p$$

RESULTS


AVX2 vs Basic Crout PLUQ

Sizes	100	300	500	1000	1200
Crout Basic (ms)	0.71	19.00	87.40	697.97	1204.63
Crout AVX2 (ms)	0.11	2.37	12.97	135.07	246.83
Speedup	6.45	8.01	6.73	5.16	4.86

Crout PLUQ Speedups using AVX2 and 12 Bits Length Prime

RESULTS

AVX2 vs FLINT



RESULTS

AVX2 vs FFLAS-FFPACK

Sizes	100	300	500	1000	1200
FFLAS-FFPACK (ms)	0.34	4.43	15.01	62.32	88.28
Crout AVX2 (ms)	0.11	2.37	12.97	135.07	246.83
Speedup	3.09	2.37	1.16	0.46	0.36

FFLAS-FFPACK Speedups using Crout PLUQ AVX2 and 12 Bits
Length Prime

RESULTS

Crout vs Basic PLUQ

Sizes	100	300	500	1000	1200
Basic PLUQ (ms)	0.69	18.70	86.57	694.06	1199.92
Crout PLUQ (ms)	0.71	19.00	87.40	697.97	1204.63
Basic PLUQ AVX2 (ms)	0.15	3.90	17.80	145.03	249.08
Crout AVX2 (ms)	0.11	2.37	12.97	135.07	246.83

Crout PLUQ vs Basic PLUQ using AVX2 and 12 Bits Length Prime