

Gaussian Elimination High-Performance Implementation

Mohamed Imad Eddine Ghodbane
May, 2024



CONTENT

- 01** PROJECT OVERVIEW
- 02** PLUQ FACTORIZATION
- 03** MODULAR ARITHMETIC USING SIMD
- 04** SIMD IMPLEMENTATION OF PLUQ
- 05** CROUT METHOD
- 06** SIMD SCALAR PRODUCT
- 07** SIMD IMPLEMENTATION OF CROUT METHOD
- 08** RESULTS

PROJECT OVERVIEW

Primary Objective

- Develop a high-performance implementation of Gaussian elimination.
- Focus on exact linear algebra over finite fields ($F_p = \mathbb{Z}/p\mathbb{Z}$) with a prime number stored on 30 bits.

Existing Libraries

- FFLAS-FFPACK, Flint, NTL.
- High performance for large matrices and small prime fields.

Areas for Improvement

- Intermediate matrix dimensions (hundreds to thousands of elements).
- Larger prime numbers.

PLUQ FACTORIZATION

$$A = PLUQ$$

Input

- Matrix A of size $m \times n$ with entries elements in the finite field F_p .

Output

- LU decomposition of A , with permutation matrices P and Q .

Key Operations

- Swap rows and columns based on pivot (pivot $\neq 0$).
- Zero out elements below the pivot (Row Reduction).

PLUQ FACTORIZATION

Column Transposition

- Let A be a 3×3 matrix:

$$A = \begin{pmatrix} a_{00} & a_{01} & a_{02} \\ a_{10} & a_{11} & a_{12} \\ a_{20} & a_{21} & a_{22} \end{pmatrix}$$

pivot = 0

- Choose the next first non-zero entry in the row.
- Perform column transposition to move the non-zero entry to the pivot position.

PLUQ FACTORIZATION

Column Transposition

$$A = \begin{pmatrix} a_{01} & 0 & a_{02} \\ a_{11} & a_{01} & a_{12} \\ a_{21} & a_{01} & a_{22} \end{pmatrix} \quad Q = \begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

pivot = 0

- Choose the next first non-zero entry in the row.
- Perform column transposition to move the non-zero entry to the pivot position.

PLUQ FACTORIZATION

Row Rotation

$$A = \begin{pmatrix} 0 & 0 & 0 \\ a_{10} & a_{11} & a_{12} \\ a_{20} & a_{21} & a_{22} \end{pmatrix}$$

pivot row = [0, 0, 0]

- Perform row rotation to move the zero row to the last row position.

PLUQ FACTORIZATION

Row Rotation

$$A = \begin{pmatrix} a_{10} & a_{11} & a_{12} \\ a_{20} & a_{21} & a_{22} \\ 0 & 0 & 0 \end{pmatrix}$$

$$P = \begin{pmatrix} 0 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \end{pmatrix}$$

pivot row = [0, 0, 0]

- Perform row rotation to move the zero row to the last row position.

PLUQ FACTORIZATION

Row Reduction

$$A = \begin{pmatrix} a_{00}^{-1} & a_{01} & a_{02} \\ a_{10} & a_{11} & a_{12} \\ a_{20} & a_{21} & a_{22} \end{pmatrix}$$

- Compute the inverse of the pivot modulo p .
- Multiply the pivot inverse by the elements of the pivot row.
- Use the updated pivot row to eliminate the elements below the pivot in the current column.

PLUQ FACTORIZATION

Row Reduction

$$\begin{pmatrix} l_{10} \\ l_{20} \end{pmatrix} = a_{00}^{-1} \cdot \begin{pmatrix} a_{10} \\ a_{20} \end{pmatrix} \qquad A = \begin{pmatrix} a_{00} & a_{01} & a_{02} \\ \textcolor{red}{l}_{10} & a_{11} & a_{12} \\ \textcolor{red}{l}_{20} & a_{21} & a_{22} \end{pmatrix}$$

- Compute the inverse of the pivot modulo p .
- Multiply the pivot inverse by the elements of the pivot row.
- Use the updated pivot row to eliminate the elements below the pivot in the current column.

PLUQ FACTORIZATION

Row Reduction

$$\begin{pmatrix} a'_{11} \\ a'_{12} \end{pmatrix} = \begin{pmatrix} a_{11} \\ a_{12} \end{pmatrix} - l_{10} \cdot \begin{pmatrix} a_{01} \\ a_{02} \end{pmatrix}$$

$$\begin{pmatrix} a'_{21} \\ a'_{22} \end{pmatrix} = \begin{pmatrix} a_{21} \\ a_{22} \end{pmatrix} - l_{20} \cdot \begin{pmatrix} a_{01} \\ a_{02} \end{pmatrix}$$

$$A = \begin{pmatrix} a_{00} & a_{01} & a_{02} \\ l_{10} & a'_{11} & a'_{12} \\ l_{20} & a'_{21} & a'_{22} \end{pmatrix}$$

- Compute the inverse of the pivot modulo p .
- Multiply the pivot inverse by the elements of the pivot row.
- Use the updated pivot row to eliminate the elements below the pivot in the current column.

PLUQ FACTORIZATION

Output

$$P = \begin{pmatrix} p_1 \\ p_2 \\ p_3 \end{pmatrix}$$

$$A = \begin{pmatrix} u_{00} & u_{01} & u_{02} \\ l_{10} & u_{11} & u_{12} \\ l_{20} & l_{21} & u_{22} \end{pmatrix}$$

$$Q = \begin{pmatrix} q_1 \\ q_2 \\ q_3 \end{pmatrix}$$

PLUQ FACTORIZATION

One Iteration

```
/*  
    Find Pivot  
*/  
  
// Row Reduction  
for (int k = matrixRank + 1; k < m; k++) {  
    A->data[k * n + matrixRank] = mult(  
        A->data[k * n + matrixRank], inv, p);  
    for (int j = matrixRank + 1; j < n; j++)  
        A->data[k * n + j] = sub(  
            A->data[k * n + j],  
            mult(A->data[k * n + matrixRank],  
                A->data[matrixRank * n + j], p), p);  
}
```

MODULAR ARITHMETIC USING SIMD

Subtraction

```
int sub(int a, int b, int p) {  
    int r = a - b;  
    return r < 0 ? r + p : r;  
}
```

Multiplication

```
int mult(int a, int b, int p) {  
    long long r = (long long)a * b;  
    return r % p;  
}
```

MODULAR ARITHMETIC USING SIMD

Subtraction

```
__m128i res = _mm_sub_epi32(a, b);  
__m128i mask = _mm_cmplt_epi32(res,  
    _mm_setzero_si128());  
__m128i ar = _mm_add_epi32(res, p);  
res = _mm_blendv_epi8(res, ar, mask);
```

a_1	a_2	a_3	a_4
-------	-------	-------	-------

-

b_1	b_2	b_3	b_4
-------	-------	-------	-------


=

$a_1 - b_1$	$a_2 - b_2$	$a_3 - b_3$	$a_4 - b_4$
-------------	-------------	-------------	-------------

MODULAR ARITHMETIC USING SIMD

Subtraction

```
__m128i res = _mm_sub_epi32(a, b);  
__m128i mask = _mm_cmplt_epi32(res,  
    _mm_setzero_si128());  
__m128i ar = _mm_add_epi32(res, p);  
res = _mm_blendv_epi8(res, ar, mask);
```



MODULAR ARITHMETIC USING SIMD

Subtraction


```
__m128i res = _mm_sub_epi32(a, b);  
__m128i mask = _mm_cmplt_epi32(res,  
                                _mm_setzero_si128());  
__m128i ar = _mm_add_epi32(res, p);  
res = _mm_blendv_epi8(res, ar, mask);
```

$a_1 - b_1$	$a_2 - b_2$	$a_3 - b_3$	$a_4 - b_4$
+			
p	p	p	p
=			
$ar1$	$ar2$	$ar3$	$ar4$

MODULAR ARITHMETIC USING SIMD

Subtraction

```
__m128i res = _mm_sub_epi32(a, b);  
__m128i mask = _mm_cmplt_epi32(res,  
    _mm_setzero_si128());  
__m128i ar = _mm_add_epi32(res, p);  
res = _mm_blendv_epi8(res, ar, mask);
```



MODULAR ARITHMETIC USING SIMD

Multiplication

```
int mult(int a, int b, int p) {  
    long long r = (long long)a * b;  
    return r % p;  
}
```

There is no modulus
operation in AVX2 !!

Solution: compute the remainder of the division by p .

$$h = ab = cp + e$$

$$d = \frac{h}{p}$$


$$c = \lfloor d \rfloor$$

$$e = h - cp$$

MODULAR ARITHMETIC USING SIMD

Multiplication


```
__m256d h = _mm256_mul_pd(a, b);  
__m256d d = _mm256_mul_pd(h, u);  
__m256d c = _mm256_floor_pd(d);  
__m256d e = _mm256_fnmadd_pd(c, p, h);
```



MODULAR ARITHMETIC USING SIMD

Multiplication

```
__m256d h = _mm256_mul_pd(a, b);  
__m256d d = _mm256_mul_pd(h, u);  
__m256d c = _mm256_floor_pd(d);  
__m256d e = _mm256_fnmadd_pd(c, p, h);
```



MODULAR ARITHMETIC USING SIMD

Multiplication

```
__m256d h = _mm256_mul_pd(a, b);  
__m256d d = _mm256_mul_pd(h, u);  
__m256d c = _mm256_floor_pd(d);  
__m256d e = _mm256_fnmadd_pd(c, p, h);
```

$$\begin{bmatrix} \lfloor d_1 \rfloor & \lfloor d_2 \rfloor & \lfloor d_3 \rfloor & \lfloor d_4 \rfloor \end{bmatrix} = \begin{bmatrix} c_1 & c_2 & c_3 & c_4 \end{bmatrix}$$

MODULAR ARITHMETIC USING SIMD

Multiplication

```
__m256d h = _mm256_mul_pd(a, b);  
__m256d d = _mm256_mul_pd(h, u);  
__m256d c = _mm256_floor_pd(d);  
__m256d e = _mm256_fnmadd_pd(c, p, h);
```

$$\begin{array}{|c|c|c|c|} \hline h_1 & h_2 & h_3 & h_4 \\ \hline \end{array} - \begin{array}{|c|c|c|c|} \hline c_1p & c_2p & c_3p & c_4p \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline e_1 & e_2 & e_3 & e_4 \\ \hline \end{array}$$

MODULAR ARITHMETIC USING SIMD

Multiplication

```
__m256d mul_mod_p(__m256d a, __m256d b, __m256d u, __m256d p) {  
    __m256d h = _mm256_mul_pd(a, b);  
    __m256d d = _mm256_mul_pd(h, u);  
    __m256d c = _mm256_floor_pd(d);  
    __m256d e = _mm256_fnmadd_pd(c, p, h);  
    return e;  
}
```

The Product exceeds the 52-bit length allocated for the mantissa in double floating-point representation. This can lead to a loss of precision in the final result !!

MODULAR ARITHMETIC USING SIMD

Multiplication

```
__m256d mul_mod_p(__m256d a, __m256d b, __m256d u, __m256d p) {  
    __m256d h = _mm256_mul_pd(a, b);  
    __m256d l = _mm256_fmsub_pd(x, y, h);  
    __m256d d = _mm256_mul_pd(h, u);  
    __m256d c = _mm256_floor_pd(d);  
    __m256d b = _mm256_fnmadd_pd(c, p, h);  
    __m256d e = _mm256_add_pd(b, l);  
    __m256d t = _mm256_sub_pd(e, p);  
    e = _mm256_blendv_pd(t, e, t);  
    t = _mm256_add_pd(e, p);  
    return _mm256_blendv_pd(e, t, e);  
}
```

SIMD IMPLEMENTATION OF PLUQ

```
rows_elimination_avx2(int *A_data, int n, int matrixRank, int c, int p ,
__m256d vp, __m256d vu , __m128i vp_128, int k) {
    __m256d vc = _mm256_set1_pd(c);
    __m256d tmp;
    int i;
    for (i = matrixRank + 1; i + 3 < n; i += 4) {
        __m128i v1 = _mm_loadu_si128((__m128i *)&A_data[matrixRank*n+i]);
        __m128i v2 = _mm_loadu_si128((__m128i *)&A_data[k*n+i]);
        __m256d vDouble = _mm256_cvtepi32_pd(v1);
        tmp = mul_mod_p(vc, vDouble, vu, vp);
        __m128i resultInt = _mm256_cvttpd_epi32(tmp);
        __m128i result = sub_avx2(v2, resultInt, vp_128);
        _mm_storeu_si128((__m128i *)&A_data[k * n + i], result);
    }
    // loop handles elements that don't fit into chunks of 4
}
```

RESULTS


AVX2 vs Basic PLUQ

Sizes	100	300	500	1000	1200
Basic (ms)	0.69	18.70	86.57	694.06	1199.92
AVX2 (ms)	0.15	3.90	17.80	139.98	242.56
Speedup	4.60	4.79	4.86	4.95	4.95

PLUQ Speedups using AVX2 and 12 Bits Length Prime

RESULTS

AVX2 vs FLINT



CROUT METHOD

Key Operation

- Computes u_{ij} and l_{ij} using the multiplication equation $A=LU$ and the previously computed values of L and U .

$$A = \begin{pmatrix} a_{00} & a_{01} & a_{02} \\ a_{10} & a_{11} & a_{12} \\ a_{20} & a_{21} & a_{22} \end{pmatrix} = P \cdot \begin{pmatrix} 1 & 0 & 0 \\ l_{10} & 1 & 0 \\ l_{20} & l_{20} & 1 \end{pmatrix} \cdot \begin{pmatrix} u_{00} & u_{01} & u_{02} \\ 0 & u_{11} & u_{12} \\ 0 & 0 & u_{22} \end{pmatrix} \cdot Q$$

$$\begin{pmatrix} l_{10} \\ l_{20} \end{pmatrix} = \begin{pmatrix} a_{10} \times u_{00}^{-1} \\ a_{20} \times u_{00}^{-1} \end{pmatrix}$$

$$A = \begin{pmatrix} u_{00} & u_{01} & u_{02} \\ l_{10} & a_{11} & a_{12} \\ l_{20} & a_{21} & a_{22} \end{pmatrix}$$

CROUT METHOD

$$A = \begin{pmatrix} u_{00} & u_{01} & u_{02} \\ l_{10} & a_{11} & a_{12} \\ l_{20} & a_{21} & a_{22} \end{pmatrix} \quad A = \begin{pmatrix} u_{00} & u_{01} & u_{02} \\ l_{10} & u_{11} & u_{12} \\ l_{20} & a_{21} & a_{22} \end{pmatrix} \quad A = \begin{pmatrix} u_{00} & u_{01} & u_{02} \\ l_{10} & u_{11} & u_{12} \\ l_{20} & l_{21} & a_{22} \end{pmatrix}$$

$$A = \begin{pmatrix} u_{00} & u_{01} & u_{02} \\ l_{10} & u_{11} & u_{12} \\ l_{20} & l_{21} & u_{22} \end{pmatrix}$$

$$u_{ij} = a_{ij} - \sum_{k=0}^{j-1} l_{ik} \cdot u_{kj}$$

$$l_{ij} = (a_{ij} - \sum_{k=0}^{i-1} l_{ik} \cdot u_{kj}) \cdot u_{jj}^{-1}$$

SIMD SCALAR PRODUCT

Scalar Product

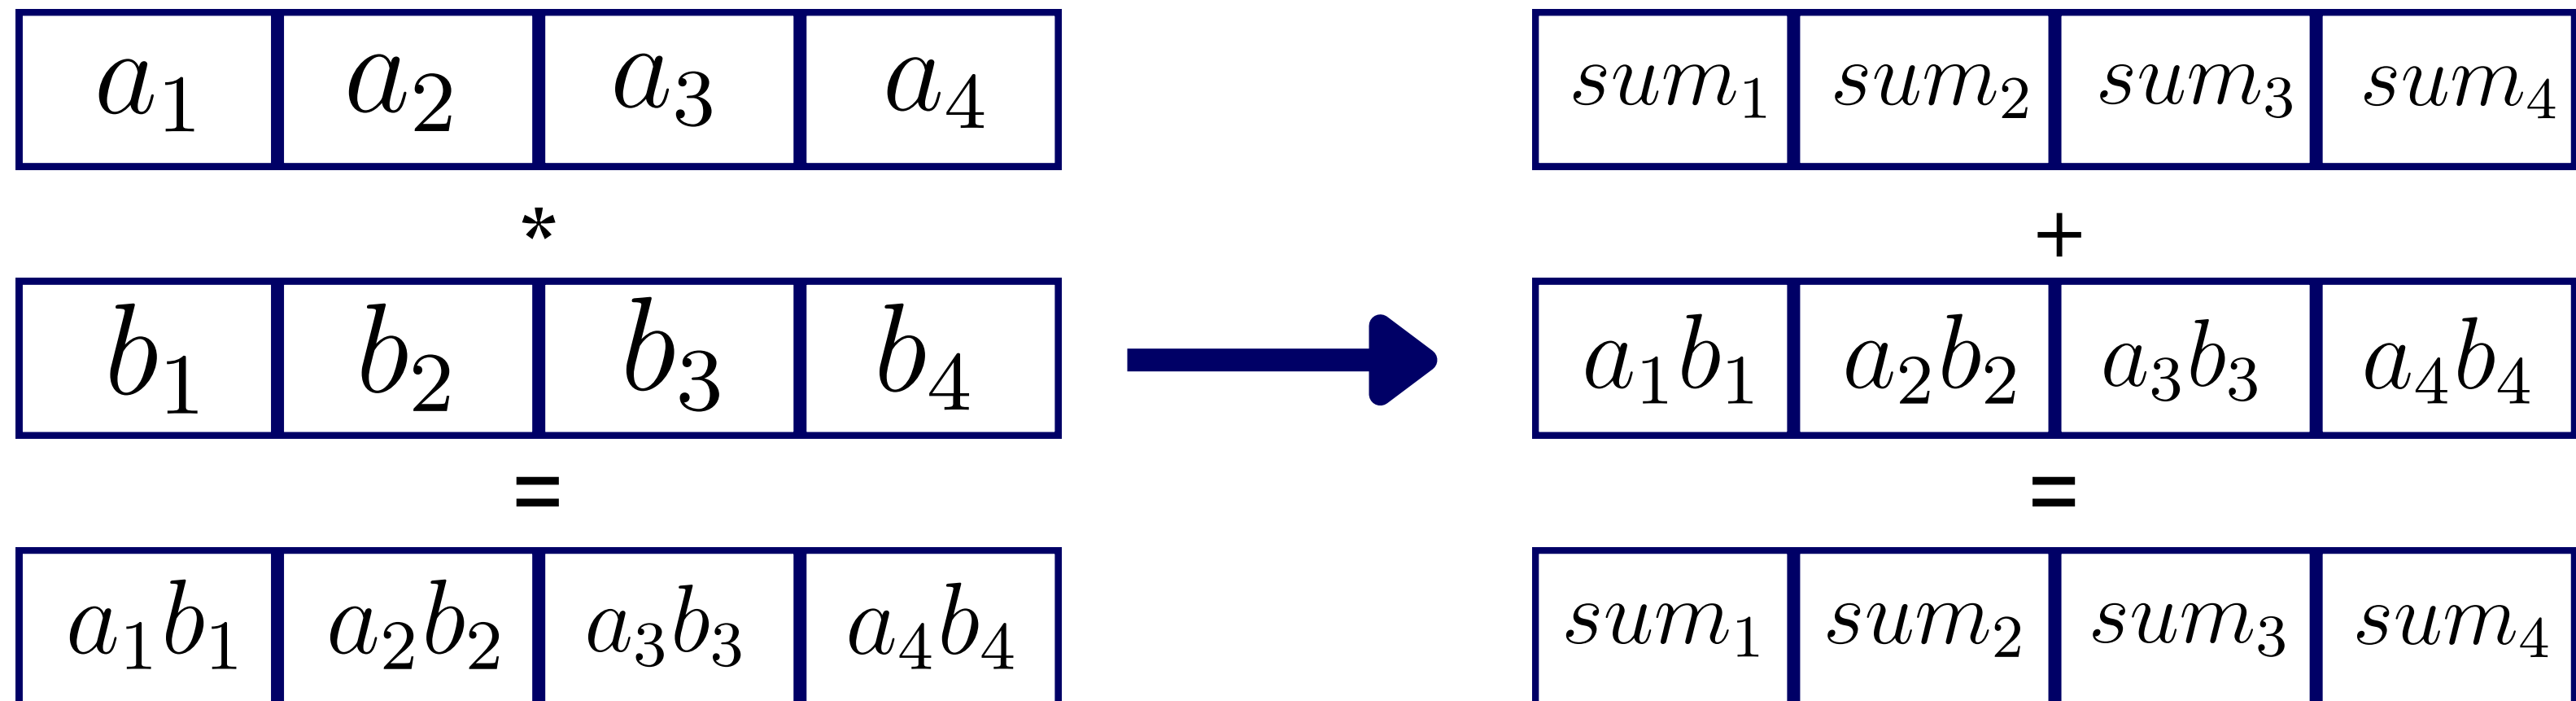
```
int sum = 0;
for (int i = 0; i < length; i++) {
    sum += a[i] * b[i];
}
return sum % p;
```

Independence of Iterations

- Each iteration of the loop computes $a[i] * b[i]$ and adds it to sum.
- The computation of $a[i] * b[i]$ is independent of the previous iterations.

SIMD SCALAR PRODUCT

Scalar Product



$$result = (sum_1 + sum_2 + sum_3 + sum_4) \bmod p$$

RESULTS


AVX2 vs Basic Crout PLUQ

Sizes	100	300	500	1000	1200
Basic (ms)	0.71	19.00	87.40	697.97	1204.63
AVX2 (ms)	0.11	2.37	12.97	135.07	246.83
Speedup	6.45	8.01	6.73	5.16	4.86

Crout PLUQ Speedups using AVX2 and 12 Bits Length Prime

RESULTS

AVX2 vs FLINT



Thank You
Questions?