

# Gaussian Elimination High-Performance Implementation

Mohamed Imad Eddine Ghodbane  
May, 2024



# CONTENT

- 01** PROJECT OVERVIEW
- 02** PLUQ FACTORIZATION
- 03** MODULAR ARITHMETIC USING SIMD
- 04** SIMD IMPLEMENTATION OF PLUQ
- 05** CROUT METHOD
- 06** SIMD SCALAR PRODUCT
- 07** SIMD IMPLEMENTATION OF CROUT METHOD
- 08** RESULTS

# PROJECT OVERVIEW

- Gaussian elimination over a finite field  $F_p$ .
- $p$  prime, less than 30 bits in length.
- Libraries include **FFLAS-FFPACK**, **Flint** known for high performance.
- Room for improvement for matrices of **intermediate dimensions**.
- **AVX2** vectorization for enhanced performance.
- Aim for superior performance compared to existing libraries.

# PLUQ FACTORIZATION

## Input:

- Matrix  $A$  of size  $m \times n$  with entries in the field  $F_p$ .

## Output:

- LU decomposition of  $A$ , with permutation matrices  $P$  and  $Q$ .

## Operations:

- Pivoting: Involves row and column rotations.
- Row Reduction: Process of reducing rows to obtain the LU decomposition.

# PLUQ FACTORIZATION

- Input

$$m = n = 3$$

$$A = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

# PLUQ FACTORIZATION

- Pivoting

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix} = [\mathbf{0}, \mathbf{0}, \mathbf{0}]$$

# PLUQ FACTORIZATION

- **Pivoting**

$$\begin{bmatrix} a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \\ a_{11} & a_{12} & a_{13} \end{bmatrix}$$

# PLUQ FACTORIZATION

- Pivoting

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix} = [\mathbf{0}, *, *]$$



# PLUQ FACTORIZATION

- **Pivoting**

$$\begin{bmatrix} a_{21} & a_{11} & a_{13} \\ a_{22} & a_{21} & a_{23} \\ a_{32} & a_{31} & a_{33} \end{bmatrix}$$

# PLUQ FACTORIZATION

- Row Reduction

$$a_{11}^{-1} \cdot \begin{pmatrix} a_{21} \\ a_{31} \end{pmatrix} = \begin{bmatrix} a_{11}^{-1} & a_{12} & a_{13} \\ l_{21} & a_{22} & a_{23} \\ l_{31} & a_{32} & a_{33} \end{bmatrix}$$

# PLUQ FACTORIZATION

- Row Reduction

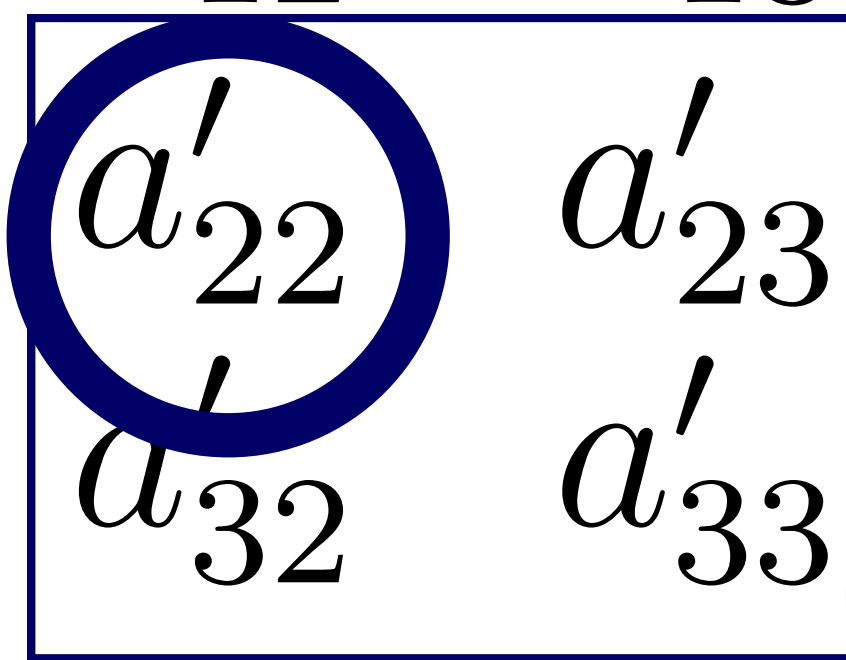
$$\begin{bmatrix} a_{11}^{-1} & a_{12} & a_{13} \\ l_{21} & a'_{22} & a'_{23} \\ l_{31} & a'_{32} & a'_{33} \end{bmatrix}$$

$$= \begin{pmatrix} a_{22} \\ a_{23} \end{pmatrix} - l_{21} \cdot \begin{pmatrix} a_{12} \\ a_{13} \end{pmatrix}$$

$$= \begin{pmatrix} a_{32} \\ a_{33} \end{pmatrix} - l_{31} \cdot \begin{pmatrix} a_{12} \\ a_{13} \end{pmatrix}$$

# PLUQ FACTORIZATION

- Row Reduction

$$\begin{bmatrix} a_{11}^{-1} & a_{12} & a_{13} \\ l_{21} & a'_{22} & a'_{23} \\ l_{31} & a'_{32} & a'_{33} \end{bmatrix}$$


# PLUQ FACTORIZATION

- Output

$$P \cdot \begin{bmatrix} u_{11} & u_{12} & u_{13} \\ l_{21} & u_{22} & u_{23} \\ l_{31} & l_{32} & u_{33} \end{bmatrix} \cdot Q$$

# PLUQ IMPLEMENTATION

```
/*  
    Pivoting  
*/  
  
// Row Reduction  
for (int k = matrixRank + 1; k < m; k++) {  
    A->data[k * n + matrixRank] = mult(  
        A->data[k * n + matrixRank], inv, p);  
    for (int j = matrixRank + 1; j < n; j++)  
        A->data[k * n + j] = sub(  
            A->data[k * n + j],  
            mult(A->data[k * n + matrixRank],  
                A->data[matrixRank * n + j], p), p);  
}
```

# MODULAR ARITHMETIC USING SIMD

- Subtraction

```
int sub(int a, int b, int p) {  
    int r = a - b;  
    return r < 0 ? r + p : r;  
}
```

# MODULAR ARITHMETIC USING SIMD

- Subtraction

`_mm_sub_epi32`

$a_1$	$a_2$	$a_3$	$a_4$
-------	-------	-------	-------

-

$b_1$	$b_2$	$b_3$	$b_4$
-------	-------	-------	-------

=

$a_1 - b_1$	$a_2 - b_2$	$a_3 - b_3$	$a_4 - b_4$
-------------	-------------	-------------	-------------



# MODULAR ARITHMETIC USING SIMD

- Subtraction

`_mm_cmplt_epi32`

$a_1 - b_1$	$a_2 - b_2$	$a_3 - b_3$	$a_4 - b_4$
-------------	-------------	-------------	-------------

<?

|   |   |   |   |
|---|---|---|---|
| 0 | 0 | 0 | 0 |
|---|---|---|---|

=

|   |   |   |   |
|---|---|---|---|
| 1 | 0 | 1 | 0 |
|---|---|---|---|

# MODULAR ARITHMETIC USING SIMD

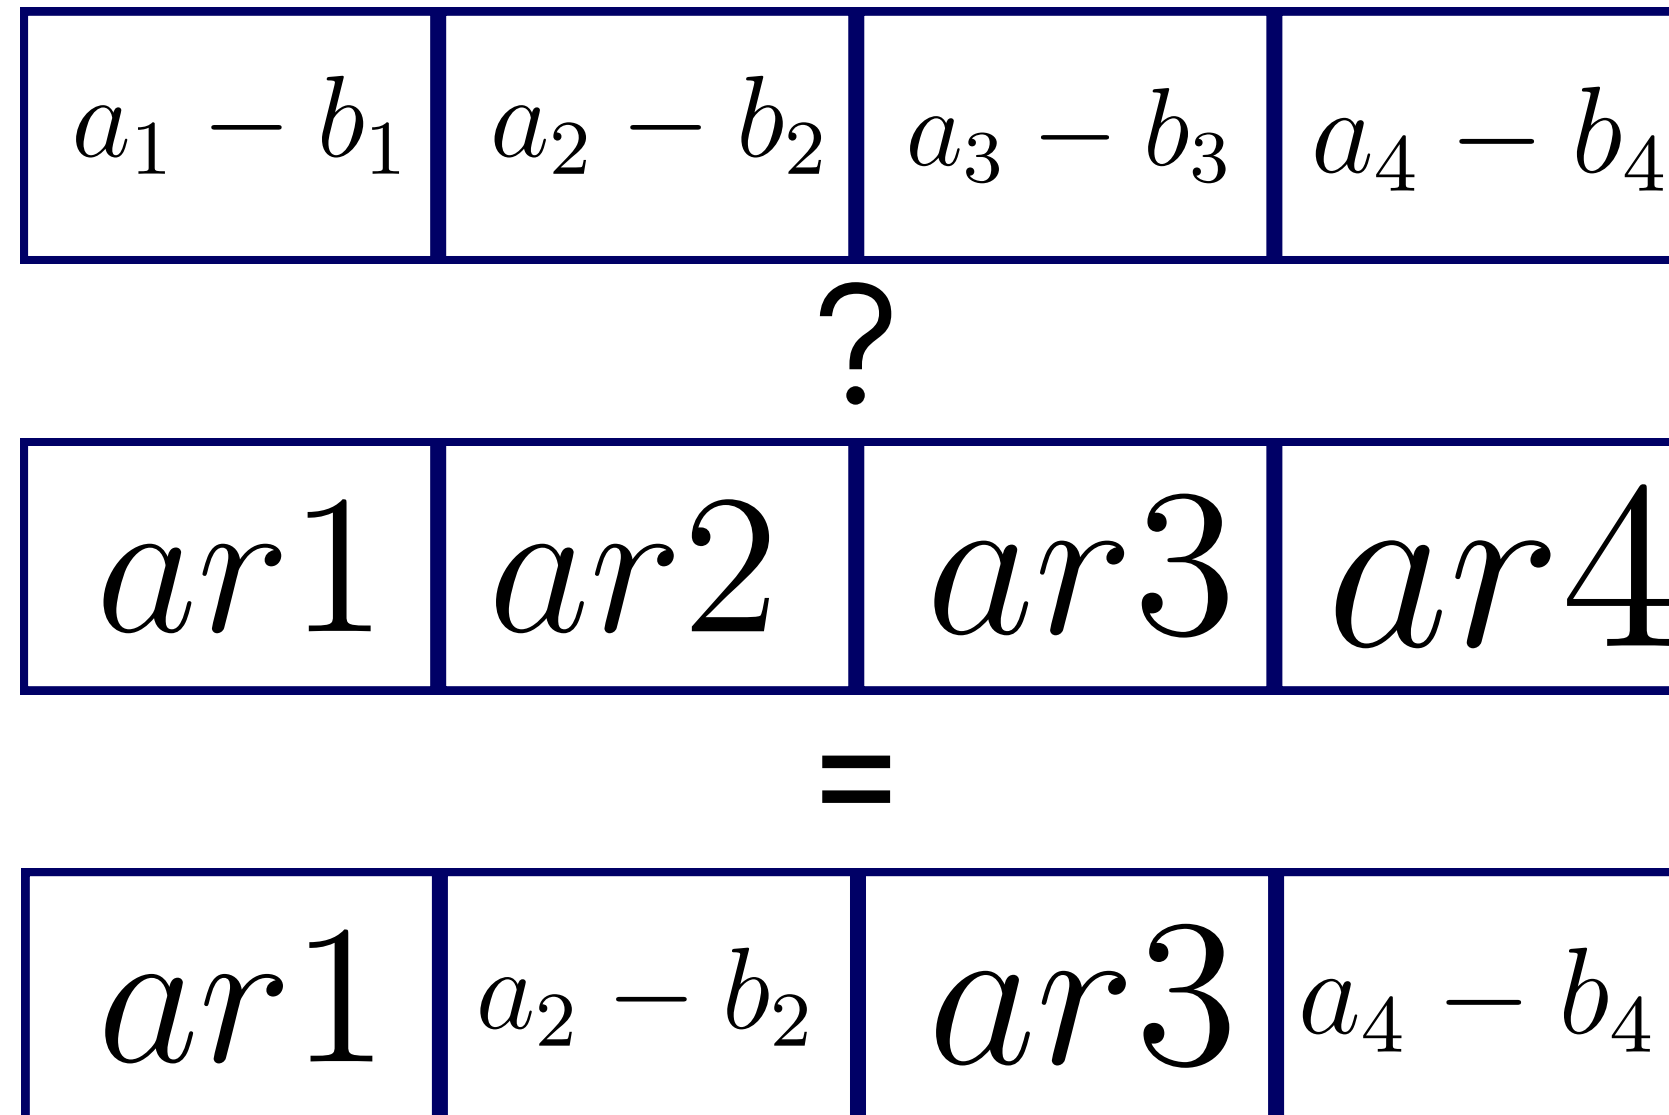
- Subtraction

`_mm_add_epi32`

|             |             |             |             |
|-------------|-------------|-------------|-------------|
| $a_1 - b_1$ | $a_2 - b_2$ | $a_3 - b_3$ | $a_4 - b_4$ |
| +           |             |             |             |
| $p$         | $p$         | $p$         | $p$         |
| =           |             |             |             |
| $ar1$       | $ar2$       | $ar3$       | $ar4$       |

# MODULAR ARITHMETIC USING SIMD

- Subtraction `_mm_blendv_epi8`



# MODULAR ARITHMETIC USING SIMD

- Subtraction

```
__m128i sub_avx2(__m128i a, __m128i b, __m128i vp) {  
    __m128i result = _mm_sub_epi32(a, b);  
    __m128i mask = _mm_cmplt_epi32(result, _mm_setzero_si128());  
    __m128i adjusted_result = _mm_add_epi32(result, vp);  
    result = _mm_blendv_epi8(result, adjusted_result, mask);  
    return result;  
}
```

# MODULAR ARITHMETIC USING SIMD

- Multiplication

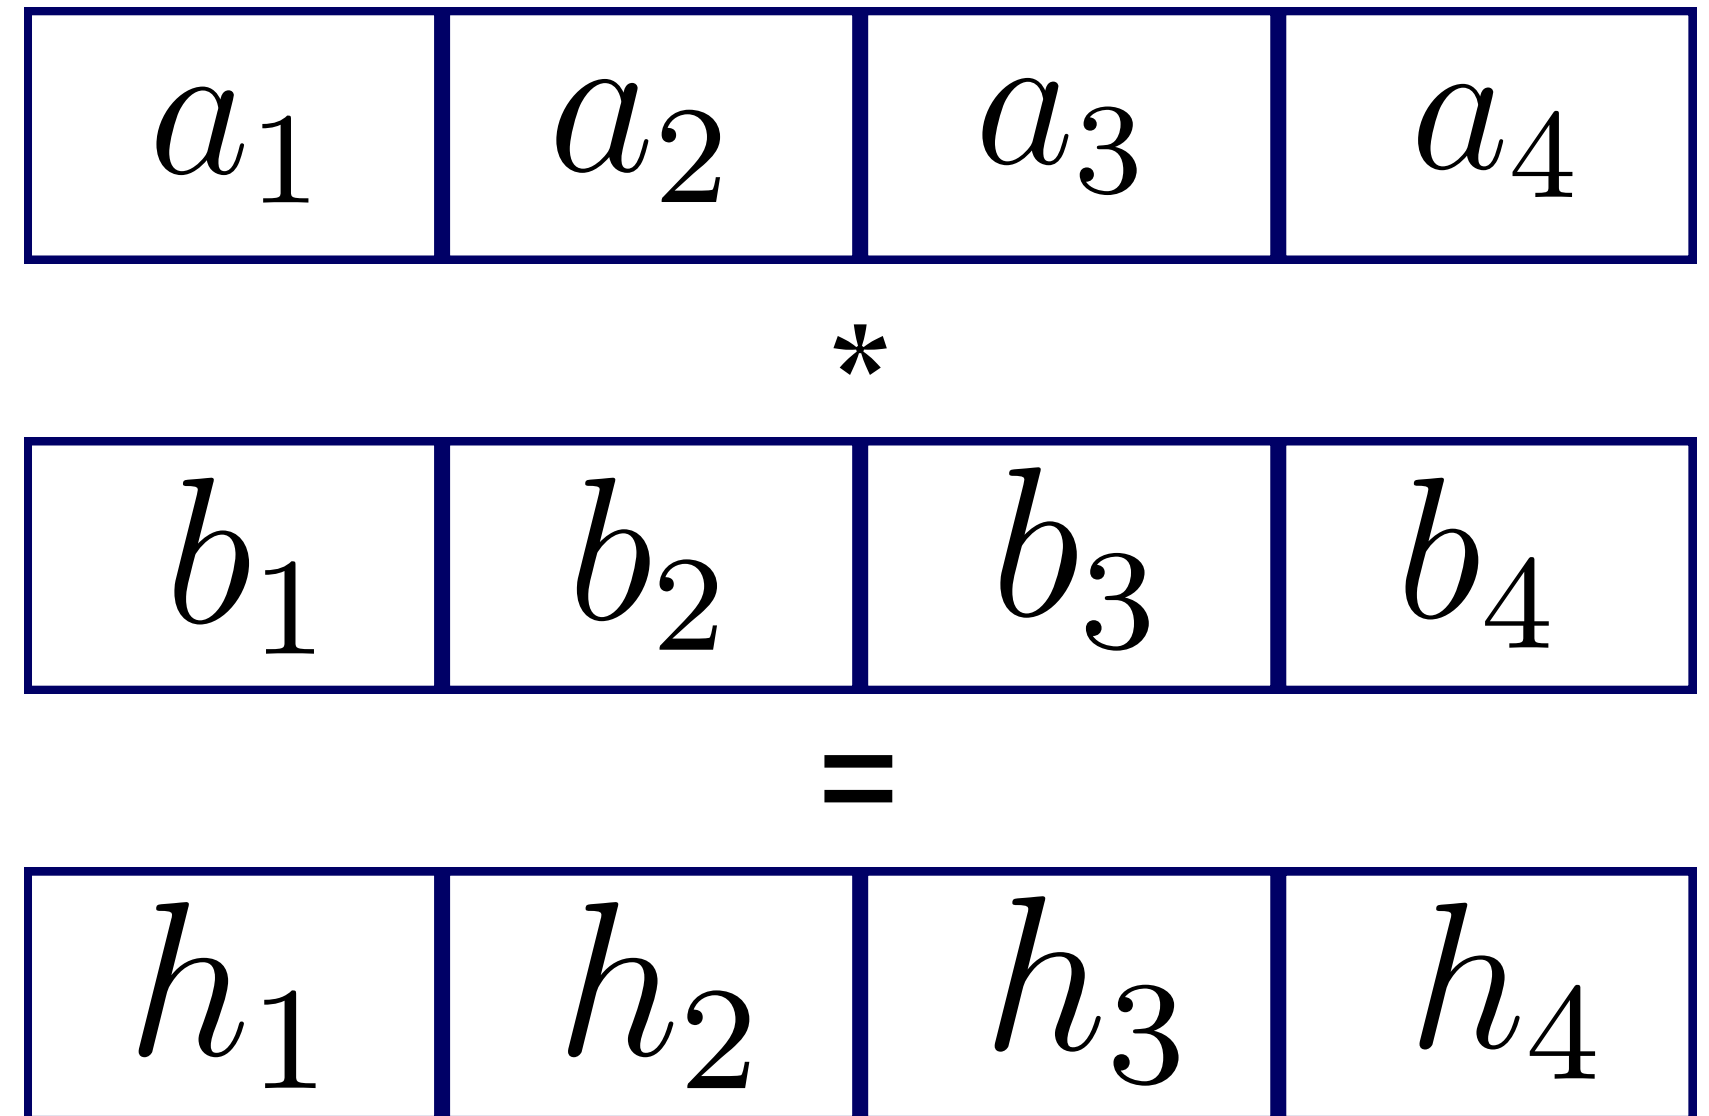
```
int mult(int a, int b, int p) {  
    long long r = (long long)a * b;  
    return r % p;  
}
```

# MODULAR ARITHMETIC USING SIMD

- Multiplication

`_mm256_mul_pd`

$$h = ab = cp + e$$



# MODULAR ARITHMETIC USING SIMD

- Multiplication

`_mm256_mul_pd`

$$h = ab = cp + e$$

$$u = \frac{1}{p}, d = hu$$

|       |       |       |       |
|-------|-------|-------|-------|
| $h_1$ | $h_2$ | $h_3$ | $h_4$ |
|-------|-------|-------|-------|

\*

|       |       |       |       |
|-------|-------|-------|-------|
| $u_1$ | $u_2$ | $u_3$ | $u_4$ |
|-------|-------|-------|-------|

=

|       |       |       |       |
|-------|-------|-------|-------|
| $d_1$ | $d_2$ | $d_3$ | $d_4$ |
|-------|-------|-------|-------|

# MODULAR ARITHMETIC USING SIMD

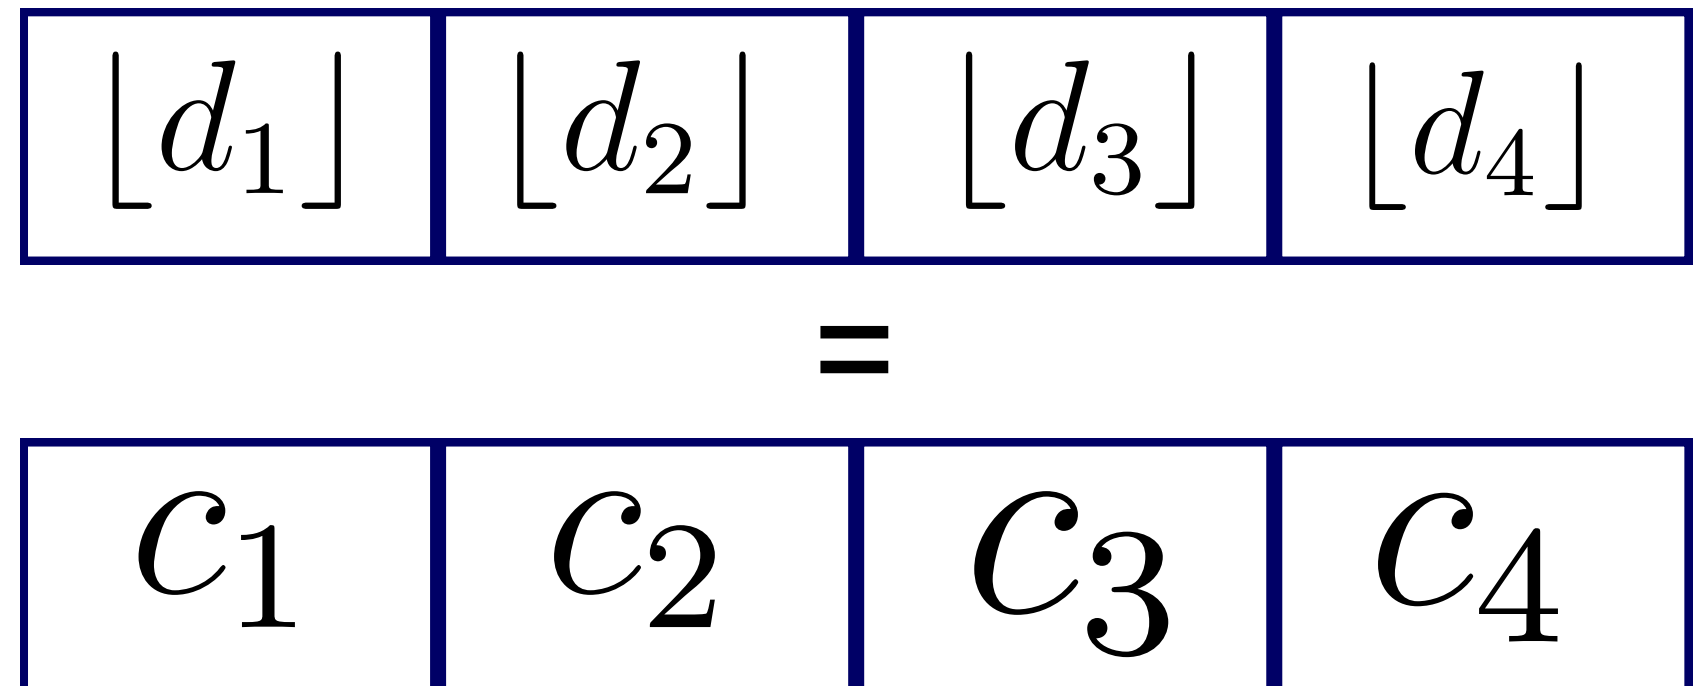
- Multiplication

`_mm256_floor_pd`

$$h = ab = cp + e$$

$$u = \frac{1}{p}, d = hu$$

$$c = \lfloor d \rfloor$$





# MODULAR ARITHMETIC USING SIMD

- Multiplication

`_mm256_fnmadd_pd`

$$h = ab = cp + e$$

$$u = \frac{1}{p}, d = hu$$

$$c = \lfloor d \rfloor$$

$$e = h - cp$$

|       |       |       |       |
|-------|-------|-------|-------|
| $h_1$ | $h_2$ | $h_3$ | $h_4$ |
|-------|-------|-------|-------|

-

|        |        |        |        |
|--------|--------|--------|--------|
| $c_1p$ | $c_2p$ | $c_3p$ | $c_4p$ |
|--------|--------|--------|--------|

=

|       |       |       |       |
|-------|-------|-------|-------|
| $e_1$ | $e_2$ | $e_3$ | $e_4$ |
|-------|-------|-------|-------|

# MODULAR ARITHMETIC USING SIMD

- Multiplication

```
__m256d mul_mod_p(__m256d a, __m256d b, __m256d u, __m256d p) {  
    __m256d h = _mm256_mul_pd(a, b);  
    __m256d d = _mm256_mul_pd(h, u);  
    __m256d c = _mm256_floor_pd(d);  
    __m256d e = _mm256_fnmadd_pd(c, p, h);  
    return e;  
}
```

The resulting product often exceeds the 52-bit length allocated for the mantissa in double floating-point representation. This can lead to a loss of precision in the final result !!

# MODULAR ARITHMETIC USING SIMD

- Multiplication

```
__m256d mul_mod_p(__m256d a, __m256d b, __m256d u, __m256d p) {  
    __m256d h = _mm256_mul_pd(a, b);  
    __m256d l = _mm256_fmsub_pd(x, y, h);  
    __m256d d = _mm256_mul_pd(h, u);  
    __m256d c = _mm256_floor_pd(d);  
    __m256d b = _mm256_fnmadd_pd(c, p, h);  
    __m256d e = _mm256_add_pd(b, l);  
    __m256d t = _mm256_sub_pd(e, p);  
    e = _mm256_blendv_pd(t, e, t);  
    t = _mm256_add_pd(e, p);  
    return _mm256_blendv_pd(e, t, e);  
}
```

# SIMD IMPLEMENTATION OF PLUQ

```
rows_elimination_avx2(int *A_data, int n, int matrixRank, int c, int p ,
__m256d vp, __m256d vu , __m128i vp_128, int k) {
    __m256d vc = _mm256_set1_pd(c);
    __m256d tmp;
    int i;
    for (i = matrixRank + 1; i + 3 < n; i += 4) {
        __m128i v1 = _mm_loadu_si128((__m128i *)&A_data[matrixRank*n+i]);
        __m128i v2 = _mm_loadu_si128((__m128i *)&A_data[k*n+i]);
        __m256d vDouble = _mm256_cvtepi32_pd(v1);
        tmp = mul_mod_p(vc, vDouble, vu, vp);
        __m128i resultInt = _mm256_cvttpd_epi32(tmp);
        __m128i result = sub_avx2(v2, resultInt, vp_128);
        _mm_storeu_si128((__m128i *)&A_data[k * n + i], result);
    } // loop handles elements that don't fit into chunks of 4
}
```

# RESULTS

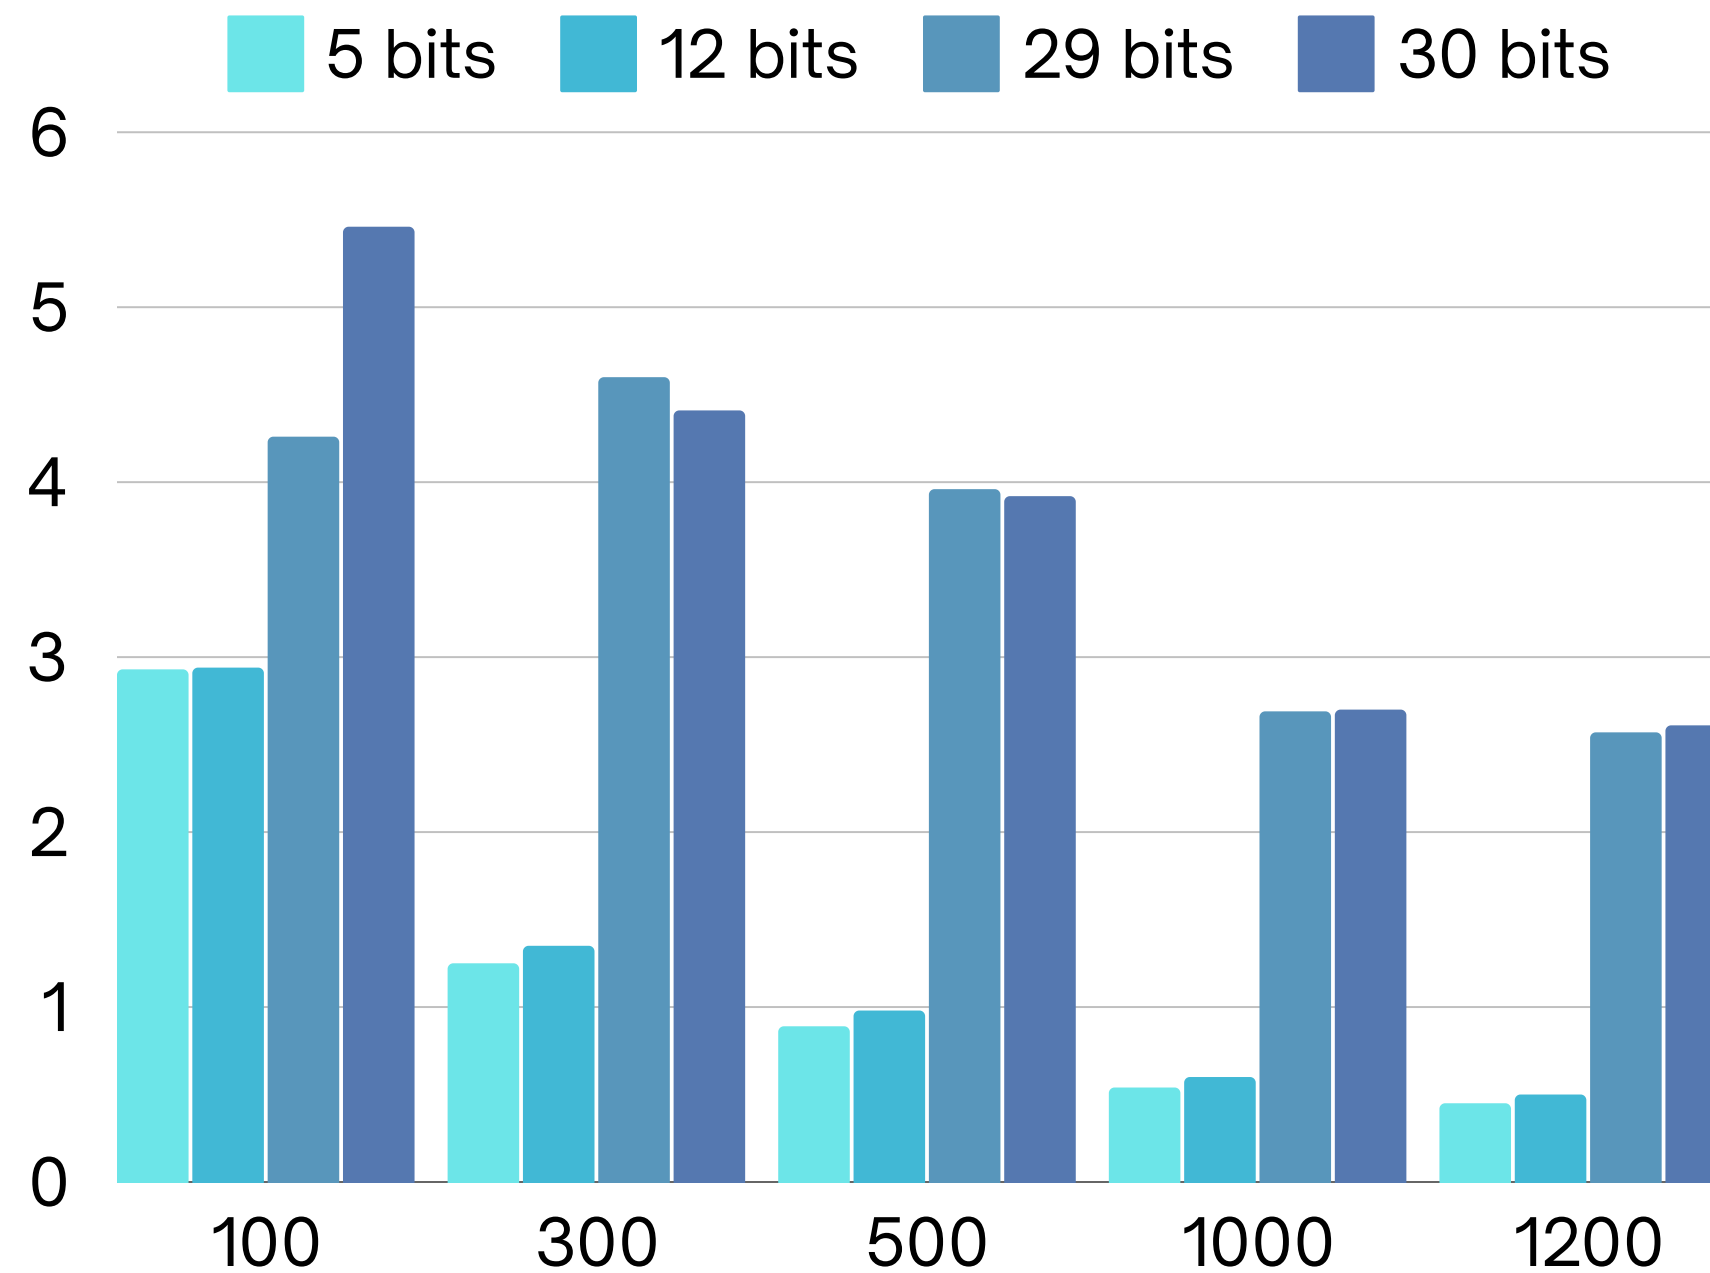
- AVX2 vs Basic PLUQ**

| Sizes      | 100  | 300   | 500   | 1000   | 1200    |
|------------|------|-------|-------|--------|---------|
| Basic (ms) | 0.69 | 18.70 | 86.57 | 694.06 | 1199.92 |
| AVX2 (ms)  | 0.15 | 3.90  | 17.80 | 139.98 | 242.56  |
| Speedup    | 4.60 | 4.79  | 4.86  | 4.95   | 4.95    |

PLUQ Speedups using AVX2 and 12 Bits Length Prime

# RESULTS

- AVX2 vs FLINT**



AMD Ryzen™ 7 PRO 7840U w/ Radeon™ 780M Graphics × 16